

Housing

August 1, 2021

```
[11]: from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

import os
import tarfile
import urllib

DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ageron/handson-ml/master/"

HOUSING_URL = DOWNLOAD_ROOT + "datasets/housing/housing.tgz"
HOUSING_PATH = os.path.join("datasets", "housing")

def fetch_housing_data(housing_url=HOUSING_URL, housing_path=HOUSING_PATH):
    if not os.path.isdir(housing_path):
        os.makedirs(housing_path)
    tgz_path = os.path.join(housing_path, "housing.tgz")
    urllib.request.urlretrieve(housing_url, tgz_path)
    housing_tgz = tarfile.open(tgz_path)
    housing_tgz.extractall(path=housing_path)
    housing_tgz.close()
    # print(os.listdir(housing_path))

import pandas as pd

def load_housing_data(housing_path=HOUSING_PATH):
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)

fetch_housing_data()
housing = load_housing_data()
housing.head()
housing.info()
```

```
[11]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	

3	-122.25	37.85	52.0	1274.0	235.0
4	-122.25	37.85	52.0	1627.0	280.0

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20640 entries, 0 to 20639
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	longitude	20640 non-null	float64
1	latitude	20640 non-null	float64
2	housing_median_age	20640 non-null	float64
3	total_rooms	20640 non-null	float64
4	total_bedrooms	20433 non-null	float64
5	population	20640 non-null	float64
6	households	20640 non-null	float64
7	median_income	20640 non-null	float64
8	median_house_value	20640 non-null	float64
9	ocean_proximity	20640 non-null	object

```
dtypes: float64(9), object(1)
```

```
memory usage: 1.6+ MB
```

```
[12]: housing.ocean_proximity.value_counts()
housing.describe()
```

```
[12]: <1H OCEAN      9136
      INLAND      6551
      NEAR OCEAN   2658
      NEAR BAY    2290
      ISLAND        5
      Name: ocean_proximity, dtype: int64
```

```
[12]:
```

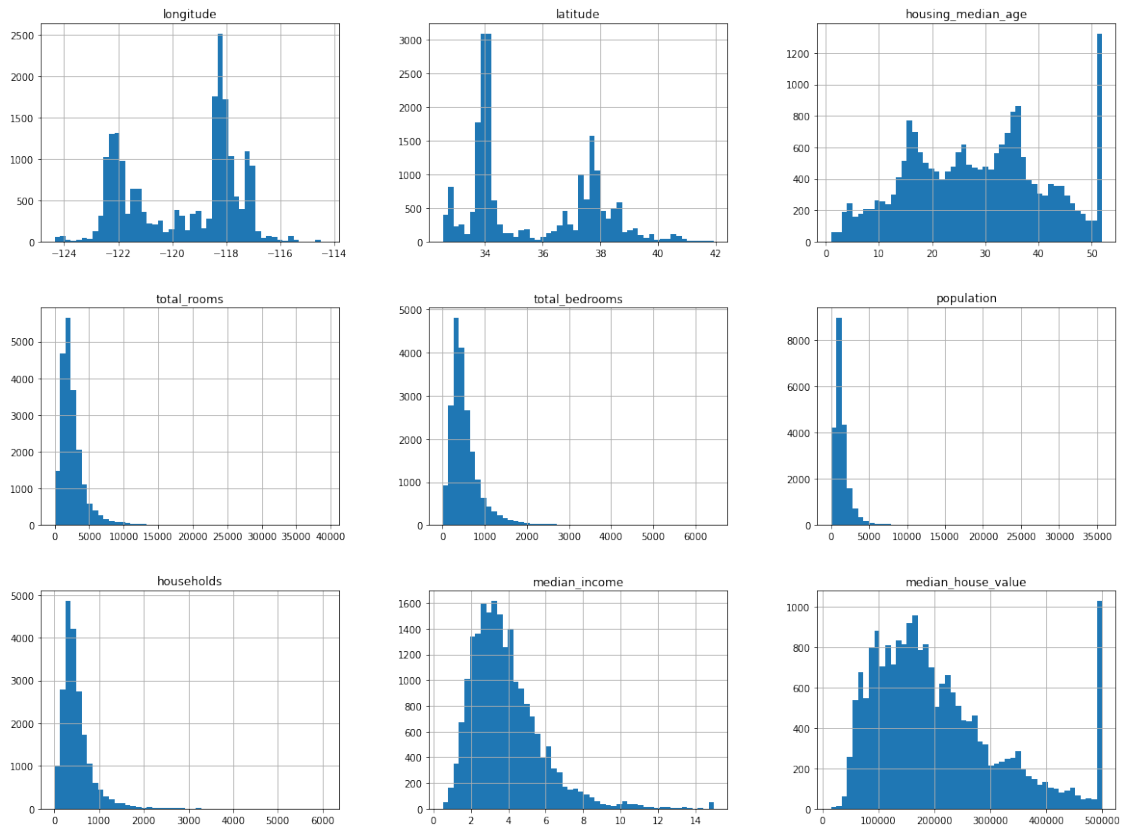
	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
[13]: %matplotlib inline
import matplotlib.pyplot as plt
housing.hist(bins=50, figsize=(20,15))
plt.show()
```

```
[13]: array([[<AxesSubplot:title={'center':'longitude'}>,
<AxesSubplot:title={'center':'latitude'}>,
<AxesSubplot:title={'center':'housing_median_age'}>],
[<AxesSubplot:title={'center':'total_rooms'}>,
<AxesSubplot:title={'center':'total_bedrooms'}>,
<AxesSubplot:title={'center':'population'}>],
[<AxesSubplot:title={'center':'households'}>,
<AxesSubplot:title={'center':'median_income'}>,
<AxesSubplot:title={'center':'median_house_value'}>]],
dtype=object)
```



```
[14]: # Create the test set
import numpy as np

# Splitting strategy 1: Shuffle. (Con: cannot handle new data.)
def split_train_set(data, test_ratio=0.2):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size = int(len(data) * test_ratio)
    test_indices = shuffled_indices[:test_set_size]
    train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices], data.iloc[test_indices]
train_set, test_set = split_train_set(housing, 0.2)
print("{} training datapoints, {} test datapoints".format(len(train_set),
    ↳len(test_set)))

import hashlib

# Splitting strategy 2: Split by identifiers (Indexed Identifier Hash vs.
    ↳Unique Feature Combination)
def test_set_check(identifier, test_ratio, hash_func):
    # Digest the identifier hash and test the last byte.
    return hash_func(np.int64(identifier)).digest()[-1] < 256 * test_ratio
```

```

def split_train_test_by_id(data, test_ratio, id_column, hash_func=hashlib.md5):
    ids = data[id_column]
    in_test_set_flags = ids.apply(lambda id: test_set_check(id, test_ratio,
↳hash_func))
    return data.loc[~in_test_set_flags], data.loc[in_test_set_flags]

housing_with_id = housing.reset_index() # Creates a new index column.
train_set, test_set = split_train_test_by_id(housing_with_id, test_ratio=0.2,
↳id_column="index")
print("{} training datapoints, {} test datapoints".format(len(train_set),
↳len(test_set)))

# Con: prone to sampling bias when multiple datapoints end up with the same
↳identifier)
housing_with_id["id"] = housing.longitude * 1000 + housing.latitude
train_set, test_set = split_train_test_by_id(housing_with_id, test_ratio=0.2,
↳id_column="id" )
print("{} training datapoints, {} test datapoints".format(len(train_set),
↳len(test_set)))

# from sklearn.model_selection import train_test_split
# train_set, test_set = train_test_split(housing, test_size=0.2,
↳random_state=42)

```

16512 training datapoints, 4128 test datapoints

16362 training datapoints, 4278 test datapoints

16267 training datapoints, 4373 test datapoints

```

[15]: # Stratified Sampling
housing["income_cat"] = np.ceil(housing.median_income / 1.5)
housing["income_cat"].where(housing["income_cat"] < 5, 5.0, inplace=True)
# housing.hist(column="income_cat", bins=50, figsize=(20,15))

from sklearn.model_selection import StratifiedShuffleSplit

split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing.income_cat):
    strat_train_set = housing.loc[train_index]
    strat_test_set = housing.loc[test_index]
strat_test_set.income_cat.value_counts() / len(strat_test_set)

# Remove incoming category column.
strat_train_set.drop("income_cat", axis=1, inplace=True)
strat_test_set.drop("income_cat", axis=1, inplace=True)


```

```
[15]: 3.0    0.350533
      2.0    0.318798
      4.0    0.176357
      5.0    0.114583
      1.0    0.039729
      Name: income_cat, dtype: float64
```

```
[16]: # Visualize and play with a sample or a copy of the training set to gain more
      ↪ insights.
housing = strat_train_set.copy()
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4,
              s=housing.population/100, label="population", figsize=(15,10.5),
              c=housing.median_house_value, cmap=plt.get_cmap("jet"),
              ↪ colorbar=True
              ) # Use more transparent dots to indicate density.
plt.legend()
```

```
[16]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```

```
[16]: <matplotlib.legend.Legend at 0x11d1400a0>
```




```
[17]: # Look for correlations.
corr_matrix = housing.corr()
corr_matrix.median_house_value.sort_values(ascending=False)

from pandas.plotting import scatter_matrix
attributes = ["median_house_value", "median_income", "total_rooms",
↳ "housing_median_age"]
scatter_matrix(housing[attributes], figsize=(12,8))
```


```
[17]: median_house_value    1.000000
median_income             0.687160
total_rooms              0.135097
housing_median_age       0.114110
households               0.064506
total_bedrooms           0.047689
population              -0.026920
longitude               -0.047432
latitude                -0.142724
Name: median_house_value, dtype: float64
```

```
[17]: array([[<AxesSubplot:xlabel='median_house_value', ylabel='median_house_value'>,
<AxesSubplot:xlabel='median_income', ylabel='median_house_value'>,
<AxesSubplot:xlabel='total_rooms', ylabel='median_house_value'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='median_house_value'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='median_income'>,
<AxesSubplot:xlabel='median_income', ylabel='median_income'>,
<AxesSubplot:xlabel='total_rooms', ylabel='median_income'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='median_income'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='total_rooms'>,
<AxesSubplot:xlabel='median_income', ylabel='total_rooms'>,
<AxesSubplot:xlabel='total_rooms', ylabel='total_rooms'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='total_rooms'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='median_income', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='total_rooms', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='housing_median_age',
ylabel='housing_median_age'>]],
dtype=object)
```



```
[18]: # Focus on the correlation between median_house_value vs. median_income.
housing.plot(kind="scatter", x="median_income", y="median_house_value", alpha=0.
↪1, figsize=(12, 8.4))
```

```
[18]: <AxesSubplot:xlabel='median_income', ylabel='median_house_value'>
```

```
[19]: #_
      ↳ longitude      latitude      housing_median_age      total_rooms      total_bedrooms
#_
      ↳ population      households      median_income      median_house_value      ocean_proximity
housing["rooms_per_household"] = housing["total_rooms"] / housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"] /_
      ↳housing["total_rooms"]
housing["population_per_household"] = housing["population"] /_
      ↳housing["households"]
corr_matrix = housing.corr()
corr_matrix.median_house_value.sort_values(ascending=False)
```

```
[19]: median_house_value      1.000000
      median_income          0.687160
      rooms_per_household    0.146285
      total_rooms            0.135097
      housing_median_age     0.114110
      households             0.064506
      total_bedrooms         0.047689
      population_per_household -0.021985
      population            -0.026920
      longitude              -0.047432
      latitude               -0.142724
```

```
bedrooms_per_room      -0.259984
Name: median_house_value, dtype: float64
```

```
[20]: # Prepare the data for machine learning algorithms
housing = strat_train_set.
```

```
File "<ipython-input-20-79d478df1636>", line 2
    housing = strat_train_set.
                        ^
SyntaxError: invalid syntax
```

```
[ ]:
```